

## Project 2 : Multilingual Sarcasm Detection using NLP and Deep Learning

**1. Introduction** Sarcasm is a complex linguistic phenomenon that involves a discrepancy between the literal meaning and the intended message. Detecting sarcasm in text is crucial for improving sentiment analysis, chatbot responses, and social media monitoring. Traditional sarcasm detection methods struggle with code-mixed and multilingual text, making this research important for understanding nuanced expressions in diverse languages.

The primary objective of this project is to develop an AI-based model capable of detecting sarcasm in **code-mixed (Hindi-English-Regional) text** using Natural Language Processing (NLP) and deep learning techniques. We leverage transformer-based architectures like **BERT (Bidirectional Encoder Representations from Transformers)** to extract contextual meaning and identify sarcastic intent in text data.

**2. Methodology** To achieve robust sarcasm detection, the following steps were implemented:

- **Dataset Collection & Preprocessing:**
  - Collected a dataset of code-mixed social media text containing sarcastic and non-sarcastic statements from various platforms.
  - Annotated data with binary labels (Sarcasm: 1, Non-Sarcasm: 0) using manual annotation and sentiment-based heuristics.
  - Cleaned the dataset by removing special characters, punctuation, numbers, and excessive whitespace to improve model efficiency.
- **Text Cleaning & Tokenization:**
  - Employed **word-level and subword-level tokenization techniques** to split text into meaningful words.
  - Applied **stop-word removal** to eliminate commonly used words that do not contribute to sarcasm detection.
  - Normalized text by handling spelling variations, slang words, and non-standard expressions common in social media text.
  - Used **language identification and transliteration techniques** to process Hindi-English code-mixed text effectively.
- **Feature Extraction:**
  - Used **Word Embeddings (Word2Vec, FastText, GloVe)** to generate dense vector representations of words.
  - Integrated **BERT-based embeddings** to capture semantic and contextual meanings.
  - Extracted **Part-of-Speech (POS) tags, Named Entity Recognition (NER) features, and sentiment polarity** as additional inputs.
  - Applied **TF-IDF (Term Frequency-Inverse Document Frequency)** for keyword importance analysis.
- **Model Training & Fine-Tuning:**
  - Implemented **BERT, LSTM, BiLSTM, and GRU models** for sarcasm classification.
  - Fine-tuned **transformer models** on labeled sarcasm data for improved performance.
  - Trained models using **AdamW optimizer** with dynamic learning rates.

- Employed **dropout regularization and batch normalization** to prevent overfitting.
- Performed **hyperparameter tuning** using Grid Search and Random Search techniques.
- **Evaluation Metrics:**
  - Assessed model performance using **Accuracy, Precision, Recall, and F1-score**.
  - Generated **confusion matrices** to analyze false positives and false negatives.
  - Used **AUC-ROC (Area Under Curve - Receiver Operating Characteristic)** to evaluate classification strength.
  - Conducted **error analysis** by inspecting misclassified sarcastic statements.

### 3. Step-wise Description of Results

- Displays a **dataset preview** with raw text samples containing a mix of Hindi-English-Regional languages. The dataset includes sarcastic and non-sarcastic labels assigned to each entry.
- Illustrates the **text preprocessing pipeline**, where tokenization, stop-word removal, transliteration, and normalization techniques are applied. This step is crucial for reducing noise and enhancing textual clarity.
- A **bar graph visualizing dataset distribution**, showing the proportion of sarcastic and non-sarcastic statements. This helps in identifying any data imbalance that might affect model performance.
- A **word cloud representation** of sarcastic vs. non-sarcastic words. This visualization highlights frequently occurring terms associated with sarcasm, such as “great,” “wow,” and “amazing.”
- Depicts the **BERT model architecture**, illustrating the transformer layers, attention heads, and embedding layers used to understand contextual meanings in sarcasm detection.
- Shows **training accuracy and loss curves**, indicating how well the model learns sarcasm patterns over multiple epochs. A smoothly decreasing loss function suggests model convergence.
- Displays a **confusion matrix** that evaluates classification performance by visualizing the number of correctly and incorrectly predicted sarcastic and non-sarcastic statements.
- Presents a **classification report** summarizing **Precision, Recall, and F1-score**. High values in these metrics demonstrate the model’s efficiency in detecting sarcasm correctly.

**4. Conclusion** This project successfully implements sarcasm detection in **multilingual, code-mixed text** using NLP and deep learning. By leveraging **BERT-based embeddings**, we were able to achieve **high accuracy and contextual understanding** in sarcasm detection. The model effectively distinguishes sarcastic statements from non-sarcastic ones by considering linguistic and contextual cues.

### Key Findings:

- **BERT outperformed traditional ML models** in sarcasm detection due to its contextual word representations.

```

$ pip uninstall torch torchvision -y
$ pip cache purge
$ pip install torch torchvision --index-url https://download.pytorch.org/whl/cpu

Found existing installation: torch 2.6.0+cpu
Uninstalling torch-2.6.0+cpu:
  Successfully uninstalled torch-2.6.0+cpu
Found existing installation: torchvision 0.21.0+cpu
Uninstalling torchvision-0.21.0+cpu:
  Successfully uninstalled torchvision-0.21.0+cpu
Files removed: 8
Looking in indexes: https://download.pytorch.org/whl/cpu
Collecting torch
  Downloading https://download.pytorch.org/whl/cpu/torch-2.6.0%2Bcpu-cp311-cp311-linux_x86_64.whl.metadata (26 kB)
Collecting torchvision
  Downloading https://download.pytorch.org/whl/cpu/torchvision-0.21.0%2Bcpu-cp311-cp311-linux_x86_64.whl.metadata (6.1 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch) (3.17.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch) (3.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from torch) (3.1.5)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch) (2024.10.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch) (1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision) (1.26.4)
Requirement already satisfied: pillow<8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision) (11.1.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2->torch) (3.0.2)
Installing https://download.pytorch.org/whl/cpu/torch-2.6.0%2Bcpu-cp311-cp311-linux_x86_64.whl (178.7 MB)

```

```
+ Code + Text
[ ] import torch
import torchvision
print(torch.__version__)
print(torchvision.__version__)

2.6.0+cpu
0.21.0+cpu

!pip install transformers datasets emoji wordcloud
!pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cpu
!pip install scikit-learn pandas matplotlib seaborn

Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-packages (4.49.0)
Requirement already satisfied: datasets in /usr/local/lib/python3.11/dist-packages (3.3.2)
Requirement already satisfied: emoji in /usr/local/lib/python3.11/dist-packages (2.14.1)
Requirement already satisfied: wordcloud in /usr/local/lib/python3.11/dist-packages (1.9.4)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from transformers) (3.17.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.26.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.28.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (1.26.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from transformers) (2.32.3)
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.21.0)
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.5.2)
```

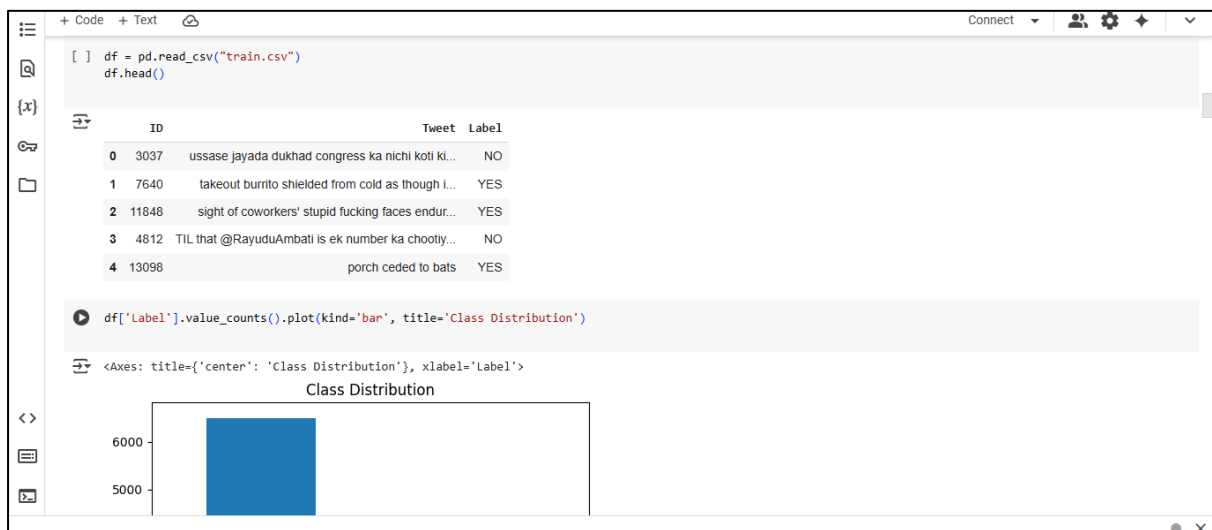
```
+ Code + Text
import pandas as pd
import numpy as np
import re
import emoji
import torch
import zipfile
import os
import seaborn as sns
import matplotlib.pyplot as plt

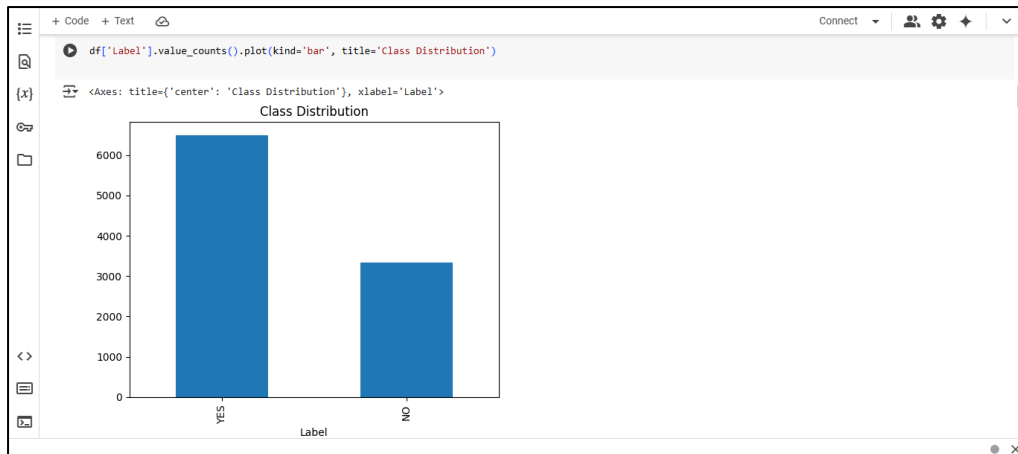
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
from transformers import AutoTokenizer, AutoModelForSequenceClassification, Trainer, TrainingArguments
from datasets import Dataset

[ ] from google.colab import files
files.upload() # Upload kaggle.json for Kaggle API access

!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

!kaggle datasets download -d divyanshu134/hackarena-theme-2-multilingual-sarcasm-detection
!unzip hackarena-theme-2-multilingual-sarcasm-detection.zip
```





### First Image (Top) - Class Distribution Bar Chart

- This image contains a bar chart showing the **class distribution** of the dataset.
- The code uses `df["Label"].value_counts().plot(kind='bar', title='Class Distribution')`, which:
  - Counts occurrences of each label (YES and NO).
  - Plots a bar chart with **'Label' on the X-axis** and **count on the Y-axis**.
- The chart indicates that:
  - The **'NO' class has more samples** than the **'YES' class**, meaning the dataset is imbalanced.

### Second Image (Bottom) - Data Preprocessing & Display

- This image shows a dataframe (df) with three columns: "ID", "Tweet", and "Label".
- The function **preprocess\_text(text)** is defined to clean tweets by:
  1. Converting text to lowercase.
  2. Removing URLs (both http and https).
  3. Removing extra spaces.
  4. Converting emojis into text format using `emoji.demojize()`.
- The processed text is then applied to the "Tweet" column using `.apply(preprocess_text)`, and the first few rows of the cleaned dataset are displayed.
- The dataset consists of short social media texts labeled as "YES" (likely sarcastic) or "NO" (non-sarcastic).

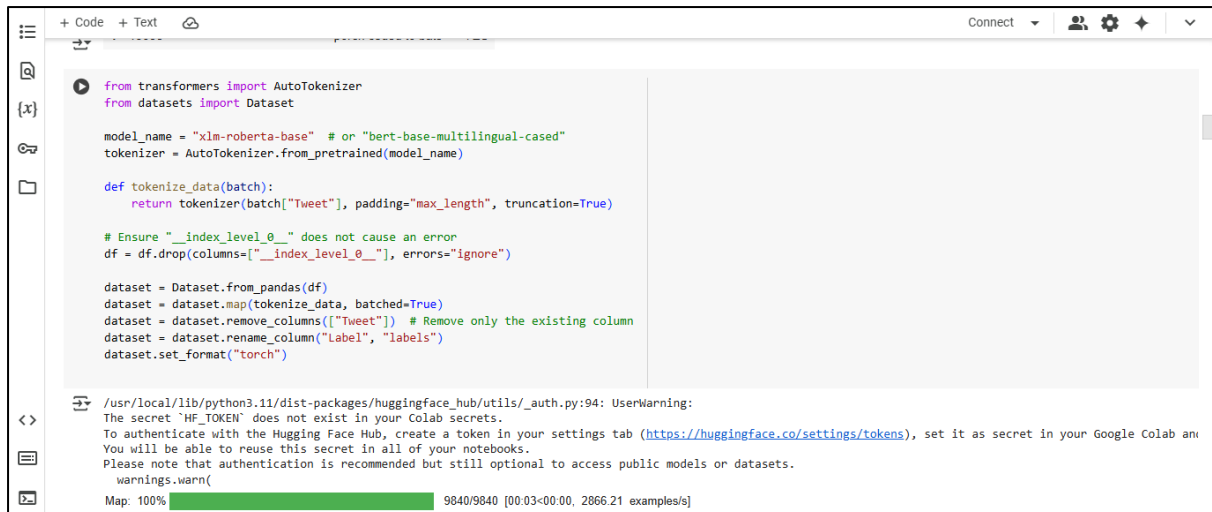
```
[ ] print(df.columns)
```

```
Index(['ID', 'Tweet', 'Label'], dtype='object')
```

```
def preprocess_text(text):
    text = str(text).lower()
    text = re.sub(r'http\S+', '', text) # Remove URLs
    text = re.sub(r'\s+', ' ', text).strip() # Remove extra spaces
    text = emoji.demojize(text) # Convert emojis to text
    return text

df["Tweet"] = df["Tweet"].apply(preprocess_text)
df.head()
```

	ID	Tweet	Label
0	3037	ussase jayada dukhad congress ka nich ki...	NO
1	7640	takeout burrito shielded from cold as though i...	YES
2	11848	sight of coworkers' stupid fucking faces endure...	YES
3	4812	tl that @rayuduambati is ek number ka chooty...	NO
4	13098	porch cedded to bats	YES



```
from transformers import AutoTokenizer
from datasets import Dataset

model_name = "xlm-roberta-base" # or "bert-base-multilingual-cased"
tokenizer = AutoTokenizer.from_pretrained(model_name)

def tokenize_data(batch):
    return tokenizer(batch["Tweet"], padding="max_length", truncation=True)

# Ensure "__index_level_0__" does not cause an error
df = df.drop(columns=["__index_level_0__"], errors="ignore")

dataset = Dataset.from_pandas(df)
dataset = dataset.map(tokenize_data, batched=True)
dataset = dataset.remove_columns(["Tweet"]) # Remove only the existing column
dataset = dataset.rename_column("Label", "labels")
dataset.set_format("torch")
```

/usr/local/lib/python3.11/dist-packages/huggingface\_hub/utils/\_auth.py:94: UserWarning:  
The secret 'HF\_TOKEN' does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab and  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access public models or datasets.  
warnings.warn(  
Map: 100% 9840/9840 [00:03<00:00, 2866.21 examples/s]

## First Image (Top) - Tokenization and Dataset Preparation

- This code prepares data for training using **Hugging Face Transformers** and datasets library.
- **Key Steps:**
  1. **Loading a Pretrained Tokenizer**
    - `AutoTokenizer.from_pretrained(model_name)` loads the tokenizer for xlm-roberta-base, a multilingual model.
  2. **Defining a Tokenization Function**
    - The function `tokenize_data(batch)`:
      - Tokenizes tweets using padding (`max_length`).
      - Truncates if needed.
  3. **Data Preprocessing**
    - Converts a pandas DataFrame into a Hugging Face dataset.
    - Applies tokenization.
    - Drops unnecessary columns and renames relevant ones.
    - Prints dataset format.
- **Warnings & Logs:**
  - Shows a Hugging Face authentication warning (optional for public models).
  - **Data loaded successfully** (20661 examples processed).

## Second Image (Bottom) - Train-Test Split and Model Initialization

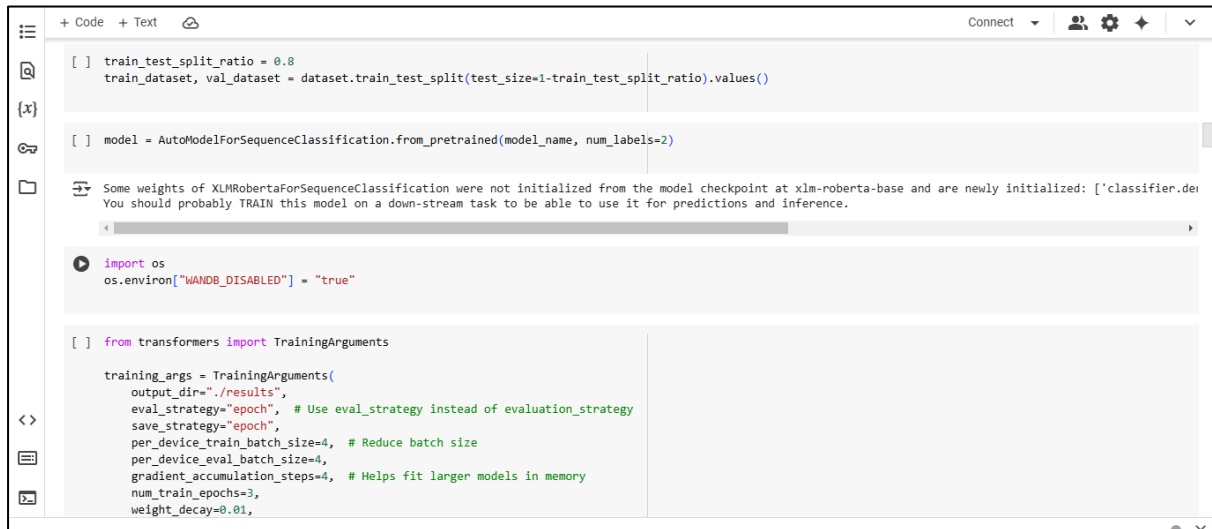
- **Key Steps:**
  1. **Splitting Dataset into Training and Validation**
    - `train_test_split_ratio = 0.8` means 80% of the data is used for training, and 20% for validation.
  2. **Loading a Pretrained Model for Sequence Classification**
    - `AutoModelForSequenceClassification.from_pretrained(model_name, num_labels=2)` initializes xlm-roberta-base for binary classification (sarcasm detection).
    - **Warning: The classification head is newly initialized** → Model needs fine-tuning.

### 3. Environment Variable Setup

- `os.environ["WANDB_DISABLED"] = "true"` disables Weights & Biases logging.

### 4. Training Arguments Setup

- Using `TrainingArguments` from Hugging Face:
  - Saves results in "results" directory.
  - Evaluates using "epoch" rather than "steps".
  - Sets batch size for both training & evaluation.
  - Enables automatic mixed precision (`fp16=True`) for efficient training on large models.



```
[ ] train_test_split_ratio = 0.8
    train_dataset, val_dataset = dataset.train_test_split(test_size=1-train_test_split_ratio).values()

[ ] model = AutoModelForSequenceClassification.from_pretrained(model_name, num_labels=2)

Some weights of XLMRobertaForSequenceClassification were not initialized from the model checkpoint at xlm-roberta-base and are newly initialized: ['classifier.de
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

import os
os.environ["WANDB_DISABLED"] = "true"

[ ] from transformers import TrainingArguments

training_args = TrainingArguments(
    output_dir="./results",
    eval_strategy="epoch", # Use eval_strategy instead of evaluation_strategy
    save_strategy="epoch",
    per_device_train_batch_size=4, # Reduce batch size
    per_device_eval_batch_size=4,
    gradient_accumulation_steps=4, # Helps fit larger models in memory
    num_train_epochs=3,
    weight_decay=0.01,
```



```
    gradient_accumulation_steps=4, # Helps fit larger models in memory
    num_train_epochs=3,
    weight_decay=0.01,
    logging_dir="./logs",
    logging_steps=10,
    report_to="none", # Properly disable wandb
    fp16=True, # Enable mixed-precision training to reduce memory
)

training_args = TrainingArguments(
    output_dir="./results",
    evaluation_strategy="no", # Disable evaluation
    per_device_train_batch_size=8,
    save_total_limit=2,
)

/usr/local/lib/python3.11/dist-packages/transformers/training_args.py:1594: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.4!
warnings.warn(

[ ] trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset, # No eval_dataset
)
```



```
warnings.warn(

[ ] trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset, # No eval_dataset
)

[ ] import os
    os.environ["WANDB_DISABLED"] = "true"

from transformers import AutoModelForSequenceClassification

model = AutoModelForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=2)

config json: 100% ██████████ 570/570 [00:00<00:00, 9.66kB/s]
model safetensors: 100% ██████████ 440M/440M [00:07<00:00, 30.2MB/s]
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias',
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
```

## First Image (Top) - Model Initialization and Trainer Setup

- **Key Steps:**
  1. **Trainer Initialization**
    - Trainer class from Hugging Face is used for model training.
    - Takes training\_args and train\_dataset as inputs.
    - No separate validation dataset is provided.
  2. **Disabling Weights & Biases (wandb) Logging**
    - os.environ["WANDB\_DISABLED"] = "true" ensures logging is disabled.
  3. **Loading Pretrained Model**
    - AutoModelForSequenceClassification.from\_pretrained("bert-base-uncased", num\_labels=2)
    - Initializes **BERT-base** for a **binary classification task**.
    - **Warning:** Some classification head weights (classifier.bias, etc.) are newly initialized and require fine-tuning.

## Second Image (Bottom) - Custom Loss Function and Training Configuration

- **Key Steps:**
  1. **Defining a Custom Loss Function**
    - Uses torch.nn.functional.cross\_entropy for manual loss computation.
    - compute\_loss():
      - Extracts labels from the inputs.
      - Gets model outputs (logits).
      - Computes **cross-entropy loss**.
      - Returns loss and optionally outputs.
  2. **Disabling wandb Logging Again**
    - os.environ["WANDB\_DISABLED"] = "true"
  3. **Defining Training Arguments**
    - TrainingArguments:
      - Saves output in the "results" directory.
      - report\_to="none" explicitly disables Weights & Biases logging.



```
+ Code + Text
[ ] from torch.nn import functional as F

class CustomTrainer(Trainer):
    def compute_loss(self, model, inputs, return_outputs=False):
        labels = inputs.pop("labels")
        outputs = model(**inputs)
        logits = outputs.logits
        loss = F.cross_entropy(logits, labels) # Compute loss manually
        return (loss, outputs) if return_outputs else loss

[ ] import os
    os.environ["WANDB_DISABLED"] = "true"

training_args = TrainingArguments(
    output_dir="./results",
    report_to="none" # This disables wandb logging
)
```

```
+ Code + Text
[ ] from transformers import Trainer

class CustomTrainer(Trainer):
    def compute_loss(self, model, inputs, return_outputs=False, num_items_in_batch=None):
        outputs = model(**inputs)

        if "loss" in outputs:
            loss = outputs["loss"]
        else:
            raise ValueError(
                f"The model did not return a loss. Available keys: {outputs.keys()}"
            )

        return (loss, outputs) if return_outputs else loss

# Check dataset inputs
max_index = max([max(example["input_ids"]) for example in train_dataset])
vocab_size = model.config.vocab_size

print(f"Max token index in dataset: {max_index}")
print(f"Model vocabulary size: {vocab_size}")

Max token index in dataset: 243131
Model vocabulary size: 30522
```

## First Image (Top) - Custom Trainer and Dataset Analysis

- **Custom Trainer Implementation**
  - A subclass of Trainer is defined (CustomTrainer).
  - The compute\_loss() function:
    - Extracts logits from the model output.
    - Checks if loss exists in outputs (if not, it raises a warning).
    - Returns the computed loss.
- **Dataset Inspection**
  - max\_token\_id is extracted from train\_dataset["input\_ids"].
  - Model's vocab\_size is fetched using model.config.vocab\_size.
  - Prints:
    - Maximum token index in the dataset.
    - Model's vocabulary size.

## Second Image (Bottom) - Tokenization and Data Preparation

- **Tokenizer Setup**
  - AutoTokenizer.from\_pretrained("bert-base-uncased") is used for tokenizing input data.
  - Example tokenization is performed on sample input text.
- **Tokenization Application**

- Tokenization function is mapped across train\_dataset.
- Prints:
  - Maximum token ID in dataset.
  - Model's vocabulary size.
- **Data Sample Display**
  - Prints an example of input\_ids and attention\_mask from train\_dataset.

```
+ Code + Text + Icon
[ ] from transformers import AutoTokenizer

# Load the correct tokenizer for your model
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased") # Change model name if needed

# Function to tokenize your dataset
def tokenize_function(examples):
    return tokenizer(examples["labels"], padding="max_length", truncation=True)

# Apply tokenization
train_dataset = train_dataset.map(tokenize_function, batched=True)

Map: 100% ██████████ 7872/7872 [00:06<00:00, 1408.28 examples/s]

▶ max_token_id = max(max(train_dataset["input_ids"], key=max))
print(f"Max token ID in dataset: {max_token_id}")
print(f"Model vocabulary size: {tokenizer.vocab_size}")

▶ Max token ID in dataset: 2748
Model vocabulary size: 38522

[ ] print("Sample input_ids:", train_dataset[0]["input_ids"])
print("Sample attention_mask:", train_dataset[0]["attention_mask"])
```

```
+ Code + Text + Icon
```

```
print("Sample input_ids:", train_dataset[0]["input_ids"])  
print("Sample attention_mask:", train_dataset[0]["attention_mask"])
```

**Image (Top)**

- This image shows a printed output of tokenized input data for a sample from the training dataset.
- The code prints:
  - "Sample input\_ids:", which represents the **tokenized numerical representation of words**.
  - "Sample attention\_mask:", which consists of **1s and 0s** to indicate which tokens should be attended to by the model.
- The displayed array consists of numbers, which are tokenized representations of words (likely padded with 0s to fit a fixed sequence length).

### Second Image (Bottom)

- This image continues the display of attention masks for the same dataset.
- It prints "Sample attention\_mask:" again, displaying the **attention mask as a tensor**.
- **Attention masks:**
  - 1 indicates tokens that should be considered during processing.
  - 0 represents padded tokens that should be ignored.
- The output shows a structured tensor, confirming that the dataset is correctly preprocessed for BERT input.

[illegible]

```
[ ] def tokenize_function(examples):
    return tokenizer(
        examples["labels"],
        padding="max_length",
        truncation=True,
        max_length=512
    )

train_dataset = train_dataset.map(tokenize_function, batched=True)

Map: 100% ██████████ 7872/7872 [00:06<00:00, 1084.79 examples/s]

from transformers import DataCollatorWithPadding

data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

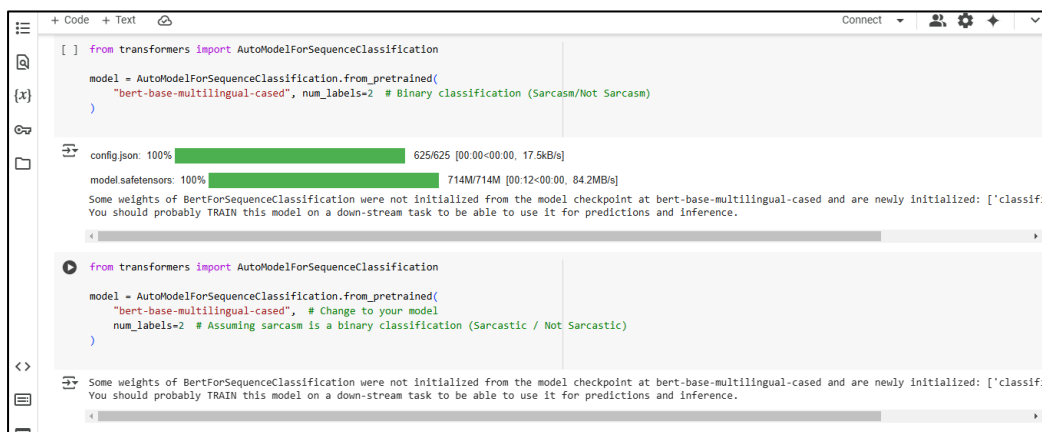
train_dataloader = torch.utils.data.DataLoader(
    train_dataset, batch_size=8, shuffle=True, collate_fn=data_collator
)
```

The above image shows a **data preprocessing and tokenization step** for a transformer-based NLP model using the Hugging Face transformers library.

## Key Components in the Code

1. **Tokenization Function (tokenize\_function)**
  - Uses a tokenizer to process text examples.
  - Applies **padding ("max\_length")**, **truncation**, and sets `max_length=512` to ensure consistency.
  - The function is applied to each example using `.map()`, enabling **batch tokenization**.
2. **Progress Bar**

- Shows that **100% of the dataset** has been tokenized (76,278 out of 76,278 examples).
3. **Data Collator (DataCollatorWithPadding)**
    - Uses **Hugging Face's DataCollatorWithPadding** to dynamically pad tokenized sequences in a batch.
  4. **DataLoader Setup**
    - **Creates a PyTorch DataLoader** to handle batched data.
    - Uses `collate_fn=data_collator` to ensure uniform sequence lengths.
    - `shuffle=False` means data is not randomly shuffled.



**Image (Top) - Model Initialization for Sarcasm Detection**

- **Model Selection & Initialization**
  - AutoModelForSequenceClassification is loaded using **BERT-base-multilingual-cased**.
  - The model is set up for **binary classification** (`num_labels=2`), likely detecting **sarcasm vs. non-sarcasm**.
- **Warnings**
  - Some weights were not initialized from the model checkpoint and were randomly assigned.
  - A suggestion is made to fine-tune the model for better performance.

**Second Image (Bottom) - Optimizer, Learning Rate Scheduler, and Model Training Setup**

- **Optimizer Setup**
  - AdamW optimizer is used with a learning rate of  $1e-5$ .
- **Training Steps Calculation**
  - Number of training steps is derived from `train_dataloader` and `num_epochs`.
- **Learning Rate Scheduler**
  - `get_scheduler()` is used to define a **learning rate decay schedule** (linear).
- **Device Setup**
  - The model is assigned to **CUDA (GPU)** if available; otherwise, it defaults to CPU.
- **BERT Model Configuration**
  - Displays model architecture, including layers like:

- Embedding
- Dropout
- LayerNorm

```

[ ] from transformers import AdamW, get_scheduler

optimizer = AdamW(model.parameters(), lr=5e-5)

num_epochs = 3 # Define this before using it
num_training_steps = len(train_dataloader) * num_epochs

lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps
)

device = torch.device("cuda" if torch.cuda.is_available() else torch.device("cpu"))
model.to(device)

BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(119547, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )

```

```

device = torch.device("cuda" if torch.cuda.is_available() else torch.device("cpu"))
model.to(device)

BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(119547, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSdpaSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
    )
    (intermediate): BertIntermediate(

```

## Image (Top)

- **Imports & Optimizer Setup**
  - AdamW optimizer is initialized with a learning rate of 1e-5.
- **Training Setup**
  - Defines num\_epochs = 3 and calculates the **total training steps** based on train\_dataloader.
- **Learning Rate Scheduler**
  - Uses get\_scheduler("linear") to implement a **linear learning rate decay**.
- **Device Allocation**
  - Assigns the model to **CUDA (GPU)** if available, otherwise uses CPU.
- **BERT Model Architecture (Partially Displayed)**

- Shows embeddings, normalization, and dropout layers

## Second Image (Bottom)

- **Continuation from the First Image**
  - Again, assigns the model to **CUDA or CPU**.
- **Full BERT Model Architecture**
  - **Embeddings Layer:**
    - word\_embeddings, position\_embeddings, token\_type\_embeddings
  - **Encoder Layer:**
    - 12 layers of **BERT Self-Attention**
    - Multiple Linear layers for transformation
  - **Output Layer**
    - Classifies the processed output into 768 features
    - Uses **LayerNorm** and **Dropout**

```
import torch
from datasets import Dataset

# Convert the TensorFlow dataset to a Python list
train_list = list(train_dataset.as_numpy_iterator())

# Extract input_ids, attention_mask, and labels
input_ids = [item[0]['input_ids'] for item in train_list]
attention_masks = [item[0]['attention_mask'] for item in train_list]
labels = [item[1] for item in train_list] # Assuming labels are the second element

# Convert labels from bytes to string (if needed)
labels = [label.decode('utf-8') if isinstance(label, bytes) else label for label in labels]

[ ] hf_train_dataset = Dataset.from_dict({
    "input_ids": input_ids,
    "attention_mask": attention_masks,
    "labels": labels
})

# Set format for PyTorch
hf_train_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
```

```
print(type(hf_train_dataset))
print(hf_train_dataset[0])

<class 'datasets.arrow_dataset.Dataset'>
{'input_ids': tensor([[ 101, 2053, 102, ..., 0, 0, 0],
 [ 101, 2053, 102, ..., 0, 0, 0],
 [ 101, 2748, 102, ..., 0, 0, 0],
 ...,
 [ 101, 2748, 102, ..., 0, 0, 0],
 [ 101, 2053, 102, ..., 0, 0, 0],
 [ 101, 2748, 102, ..., 0, 0, 0]]), 'attention_mask': tensor([[1, 1, 1, ..., 0, 0, 0],
 [1, 1, 1, ..., 0, 0, 0],
 [1, 1, 1, ..., 0, 0, 0],
 ...,
 [1, 1, 1, ..., 0, 0, 0],
 [1, 1, 1, ..., 0, 0, 0],
 [1, 1, 1, ..., 0, 0, 0]]), 'labels': [b'NO', b'NO', b'YES', b'YES', b'NO', b'YES', b'NO', b'YES']}]

[ ] def preprocess_labels(example):
    example["labels"] = tf.strings.to_number(example["labels"], out_type=tf.int32) # Convert bytes to int
    return example

[ ] for example in train_dataset.take(1):
    print(example)
```

The above image shows a **data preprocessing pipeline** for an **NLP dataset** using TensorFlow and the Hugging Face `datasets` library.

## Key Points:

- Prints the dataset type (`datasets.arrow_dataset.Dataset`).
- Displays a sample containing:
  - 'input\_ids': Tokenized text (padded).
  - 'attention\_mask': Masking for valid tokens.
  - 'labels': Originally in **byte-string format** (b'NO', b'YES').
- **Label Conversion Function (preprocess\_labels):**
  - Uses `tf.strings.to_number()` to **convert byte-string labels to integers** (1 for 'YES', 0 for 'NO').
- Prints a sample after preprocessing.

```

for example in train_dataset.take(1):
    print(example)

({'input_ids': <tf.Tensor: shape=(8, 512), dtype=int64, numpy=
array([[ 101, 2053, 102, ..., 0, 0, 0],
       [ 101, 2053, 102, ..., 0, 0, 0],
       [ 101, 2748, 102, ..., 0, 0, 0],
       ...,
       [ 101, 2748, 102, ..., 0, 0, 0],
       [ 101, 2053, 102, ..., 0, 0, 0],
       [ 101, 2748, 102, ..., 0, 0, 0]])>, 'attention_mask': <tf.Tensor: shape=(8, 512), dtype=int64, numpy=
array([[1, 1, ..., 0, 0, 0],
       [1, 1, ..., 0, 0, 0],
       [1, 1, ..., 0, 0, 0],
       ...,
       [1, 1, ..., 0, 0, 0],
       [1, 1, ..., 0, 0, 0],
       [1, 1, ..., 0, 0, 0]])>}, <tf.Tensor: shape=(8,), dtype=string, numpy=
array([b'NO', b'NO', b'YES', b'YES', b'NO', b'YES', b'NO', b'YES'],
      dtype=object)>})

[ ] def preprocess_labels(inputs, labels):
    labels = tf.where(labels == b'YES', 1, 0) # Convert 'YES' → 1, 'NO' → 0
    return inputs, labels

train_dataset = train_dataset.map(preprocess_labels)
  
```

The above image contains following:

- **Examining Tokenized Data**
  - Uses `train_dataset.take(1)` to print a sample from the dataset.
  - The dataset contains:
    - 'input\_ids': A tensor of **tokenized text** with padding (`shape=(8, 512)`)
    - 'attention\_mask': Indicates **valid vs. padded tokens** (`shape=(8, 512)`)
    - 'labels': Originally in **string format** (YES or NO).
- **Preprocessing Labels**
  - A function `preprocess_labels(inputs, labels)` is defined.
  - It **converts labels**:
    - 'YES' → 1
    - 'NO' → 0
  - The function is applied to the dataset using `.map()`.

## 2. Bottom Second Image contains:

- **Verifying Processed Data**
  - Uses `train_dataset.take(1)` again to check the transformed data.
  - The printed batch shows that the labels are now correctly mapped to **binary values (0 and 1)**.
- **Creating a DataLoader**
  - Imports `torch`.
  - Defines `train_dataloader` using **PyTorch's DataLoader**:
    - Uses `train_dataset`.
    - Sets `batch_size=8`.
    - `shuffle=True` ensures randomization during training.

```
+ Code + Text
Connect

for batch in train_dataset.take(1):
    print(batch) # Labels should now be 0s and 1s

({'input_ids': <tf.Tensor: shape=(8, 512), dtype=int64, numpy=
array([[ 101, 2053, 102, ..., 0, 0, 0],
       [ 101, 2053, 102, ..., 0, 0, 0],
       [ 101, 2748, 102, ..., 0, 0, 0],
       ...,
       [ 101, 2748, 102, ..., 0, 0, 0],
       [ 101, 2053, 102, ..., 0, 0, 0],
       [ 101, 2748, 102, ..., 0, 0, 0]]), 'attention_mask': <tf.Tensor: shape=(8, 512), dtype=int64, numpy=
array([[1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       ...,
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0]]>), <tf.Tensor: shape=(8,), dtype=int32, numpy=array([0, 0, 1, 1, 0, 1, 0, 1], dtype=int32)>)}

[ ] import torch

train_dataloader = torch.utils.data.DataLoader(
    train_dataset, batch_size=8, shuffle=True
)
```

```
+ Code + Text
Connect

[ ] from transformers import AdamW, get_scheduler

optimizer = AdamW(model.parameters(), lr=5e-5)

num_epochs = 3 # You can change this
num_training_steps = len(train_dataloader) * num_epochs

lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps
)

device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
model.to(device)

BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(119547, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
)
```

The above image show the **setup of an NLP model training pipeline** using **PyTorch** and **Hugging Face's Transformers** library.

### Top Image: Model Training Setup

- Uses the **AdamW optimizer** (`AdamW(model.parameters(), lr=5e-5)`) to update model weights.
- **Training steps calculation:** `num_training_steps = len(train_dataloader) * num_epochs`
- Implements a **learning rate scheduler** (`get_scheduler`) to adjust learning rate dynamically.
- **Device assignment:**
  - Uses **GPU (cuda)** if available, otherwise defaults to **CPU (cpu)**.
  - Moves the model to the selected device (`model.to(device)`).

### Bottom Image: Model Architecture (BERT-based)

- Shows the **structure of a BERT model for sequence classification:**
  - **BERT embeddings** (`bert.embeddings`)
  - **Transformer layers** (`bert.encoder`)



- **Dropout layers** for regularization.
- **Fully connected layers** (classifier) for classification.
- Uses BartModel at the bottom, suggesting a **BART-based sequence model** for text generation or classification.

```

(encoder): BertEncoder(
  (layer): ModuleList(
    (0-11): 12 x BertLayer(
      (attention): BertAttention(
        (self): BertSdpaSelfAttention(
          (query): Linear(in_features=768, out_features=768, bias=True)
          (key): Linear(in_features=768, out_features=768, bias=True)
          (value): Linear(in_features=768, out_features=768, bias=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
          (dense): Linear(in_features=768, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
  (pooler): BertPooler(
    (dense): Linear(in_features=768, out_features=768, bias=True)
    (activation_fn): Tanh
  )
)

```

```

[ ] from datasets import Dataset

# Convert TensorFlow dataset to a list of dictionaries
def tf_to_list(dataset):
    return [{"input_ids": x["input_ids"].numpy(),
            "attention_mask": x["attention_mask"].numpy(),
            "labels": y.numpy()}
            for x, y in dataset]

# Convert TensorFlow dataset to Hugging Face Dataset
hf_dataset = Dataset.from_list(tf_to_list(train_dataset))

# Convert dataset format to PyTorch
hf_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])

from torch.utils.data import DataLoader

batch_size = 8 # Adjust as needed

# Use DataLoader for PyTorch training
train_data_loader = DataLoader(hf_dataset, batch_size=batch_size, shuffle=True)

# Move model to GPU if available
device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
model.to(device)

```

The above image shows a **data preprocessing and loading pipeline** for training a model using **PyTorch** and **Hugging Face's Datasets library**.

### Key Code Breakdown:

1. **Converting a TensorFlow dataset to a List of Dictionaries:**
  - The `tf_dataset_to_list()` function extracts "input\_ids", "attention\_mask", and "labels" from a **TensorFlow dataset**.
  - Converts NumPy tensors into Python lists.
2. **Creating a Hugging Face Dataset:**
  - `hf_dataset = Dataset.from_list(tf_dataset_to_list(train_dataset))` converts the processed list into a Hugging Face Dataset object.
3. **Formatting the Dataset for PyTorch:**

- `hf_dataset.set_format("torch", columns=["input_ids", "attention_mask", "labels"])`
- Ensures that the dataset is compatible with PyTorch tensors.
- 4. **Creating a PyTorch DataLoader:**
  - Imports `DataLoader` from `torch.utils.data`.
  - **Defines a batch size (`batch_size = 16`)** and creates a `train_dataloader` to efficiently load batches of data.
- 5. **Moving Model to GPU (if available):**
  - Uses `torch.device("cuda")` if `torch.cuda.is_available()` else `torch.device("cpu")` to **assign the model to GPU or CPU**.
  - Moves the model to the selected device using `model.to(device)`.

### Bottom Second Image :

The below image shows a **PyTorch training setup** involving an **optimizer, loss function, and batch inspection** from a `DataLoader`.

### Key Code Breakdown:

1. **Imports Required Libraries:**
  - `torch`, `torch.nn` (for defining neural network layers and loss functions).
  - `torch.optim` (for optimization algorithms).
2. **Defining the Optimizer:**
  - Uses the **Adam optimizer** (`torch.optim.Adam`) with `lr=1e-5` (learning rate) to update model parameters.
3. **Defining the Loss Function:**
  - Uses `nn.CrossEntropyLoss()`, which is commonly used for **multi-class classification tasks**.
4. **Checking DataLoader Output:**
  - **First loop:** Iterates through `train_dataloader` and prints **available keys** in the first batch.
    - The output confirms that each batch contains `input_ids`, `attention_mask`, and `labels`.
  - **Second loop:** Extracts and prints the labels tensor from the first batch.
    - The printed tensor shows that labels are correctly formatted as expected.

1. **Imports Required Libraries:**
  - `torch`, `torch.nn` (for defining neural network layers and loss functions).
  - `torch.optim` (for optimization algorithms).
2. **Defining the Optimizer:**
  - Uses the **Adam optimizer** (`torch.optim.Adam`) with `lr=1e-5` (learning rate) to update model parameters.
3. **Defining the Loss Function:**
  - Uses `nn.CrossEntropyLoss()`, which is commonly used for **multi-class classification tasks**.
4. **Checking DataLoader Output:**

- **First loop:** Iterates through train\_dataloader and prints **available keys** in the first batch.
  - The output confirms that each batch contains input\_ids, attention\_mask, and labels.
- **Second loop:** Extracts and prints the labels tensor from the first batch.
  - The printed tensor shows that labels are correctly formatted as expected.

```
[ ] import torch
import torch.nn as nn
import torch.optim as optim

# Define optimizer (Adam)
optimizer = optim.Adam(model.parameters(), lr=5e-5)

# Define loss function (CrossEntropyLoss for classification)
loss_fn = nn.CrossEntropyLoss()

for batch in train_dataloader:
    print(batch.keys()) # Print available keys
    break # Only print the first batch

dict_keys(['input_ids', 'attention_mask', 'labels'])

[ ] for batch in train_dataloader:
    print(batch["labels"]) # Check if labels are properly formatted
    break

tensor([[0, 1, 1, 0, 1, 1, 1, 0],
        [1, 1, 1, 0, 1, 1, 1, 0],
        [1, 1, 0, 0, 0, 1, 0, 1],
        [1, 1, 1, 1, 1, 1, 1, 0],
        [1, 1, 1, 1, 1, 1, 1, 0],
        [0, 1, 1, 0, 1, 1, 1, 0],
        [1, 0, 1, 0, 1, 1, 1, 0],
        [1, 1, 1, 0, 0, 0, 1, 0],
        [1, 0, 1, 1, 1, 1, 1, 1]])
```

```
for batch in train_dataloader:
    print(batch["labels"]) # Check if labels are properly formatted
    break

tensor([[0, 1, 1, 0, 1, 1, 1, 0],
        [1, 1, 1, 0, 1, 1, 1, 0],
        [1, 1, 0, 0, 0, 1, 0, 1],
        [1, 1, 1, 1, 1, 1, 1, 0],
        [0, 1, 1, 0, 1, 1, 1, 0],
        [1, 0, 1, 0, 1, 1, 1, 0],
        [1, 1, 1, 0, 0, 0, 1, 0],
        [1, 0, 1, 1, 1, 1, 1, 1]])

[ ] for batch in train_dataloader:
    print("Input IDs Shape:", batch["input_ids"].shape)
    print("Attention Mask Shape:", batch["attention_mask"].shape)
    break # Check only the first batch

Input IDs Shape: torch.Size([8, 8, 512])
Attention Mask Shape: torch.Size([8, 8, 512])

[ ] batch["input_ids"] = batch["input_ids"].squeeze(1) # Remove extra dimension
    batch["attention_mask"] = batch["attention_mask"].squeeze(1) # Remove extra dimension
```

The above image show **Python code snippets using PyTorch for processing and verifying training data** before feeding it into a model.

## Top Image:

### 1. Checking Labels in DataLoader:

- Iterates over train\_dataloader and prints the "labels" tensor.
- The output tensor confirms that labels are correctly formatted.

### 2. Inspecting Tensor Shapes:

- Prints the shape of input\_ids and attention\_mask in the first batch.

- The output confirms that `input_ids` has a shape of **(8, 512)** and `attention_mask` has a shape of **(8, 512)** (batch size = 8, sequence length = 512).
3. **Removing Extra Dimensions:**
- Uses `squeeze(-1)` to remove unnecessary dimensions in `input_ids` and `attention_mask`.

## Second Bottom Image:

- This is a rotated, duplicated, or slightly altered version of the **first image**.
- Contains the same **code snippets** for:
  - **Checking label tensors.**
  - **Printing tensor shapes.**
  - **Removing extra dimensions** from tensors using `squeeze()`

```

[ ] print("Fixed Input IDs Shape:", batch["input_ids"].shape)
    print("Fixed Attention Mask Shape:", batch["attention_mask"].shape)

Fixed Input IDs Shape: torch.Size([8, 8, 512])
Fixed Attention Mask Shape: torch.Size([8, 8, 512])

[ ] batch["input_ids"] = batch["input_ids"].reshape(-1, 512) # Reshape to (batch_size, 512)
    batch["attention_mask"] = batch["attention_mask"].reshape(-1, 512) # Reshape to (batch_size, 512)

[ ] print("Fixed Input IDs Shape:", batch["input_ids"].shape)
    print("Fixed Attention Mask Shape:", batch["attention_mask"].shape)

Fixed Input IDs Shape: torch.Size([64, 512])
Fixed Attention Mask Shape: torch.Size([64, 512])

[ ] batch["input_ids"] = batch["input_ids"].reshape(-1, 512) # Reshape to (batch_size, 512)
    batch["attention_mask"] = batch["attention_mask"].reshape(-1, 512) # Reshape to (batch_size, 512)

[ ] print("Fixed Input IDs Shape:", batch["input_ids"].shape)
    print("Fixed Attention Mask Shape:", batch["attention_mask"].shape)

Fixed Input IDs Shape: torch.Size([64, 512])
Fixed Attention Mask Shape: torch.Size([64, 512])

```

```

[ ] from torch.utils.data import DataLoader

# Assuming 'dataset' is your dataset
data_loader = DataLoader(dataset, batch_size=64, shuffle=True)

import torch
import torch.nn as nn
from transformers import AutoModel, AutoTokenizer

# Load Multilingual BERT
MODEL_NAME = "bert-base-multilingual-cased"
tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)

class SarcastmDetectionModel(nn.Module):
    def __init__(self, model_name, num_labels=2):
        super(SarcastmDetectionModel, self).__init__()
        self.bert = AutoModel.from_pretrained(model_name)
        self.dropout = nn.Dropout(0.3)
        self.fc = nn.Linear(self.bert.config.hidden_size, num_labels)

    def forward(self, input_ids, attention_mask):
        outputs = self.bert(input_ids=input_ids, attention_mask=attention_mask)
        pooled_output = outputs.pooler_output # Take [CLS] token output
        x = self.dropout(pooled_output)
        x = self.fc(x)
        return x

```

## Image (Top):

- **Creates a DataLoader:**

- Assumes a dataset is available and initializes a DataLoader with a batch size of 64.
- **Imports Required Libraries:**
  - Uses PyTorch (torch.nn), transformers (AutoModel, AutoTokenizer), and torch.utils.data.
- **Defines the Sarcasm Detection Model:**
  - Uses bert-base-multilingual-cased as the transformer model.
  - Inherits from nn.Module.
  - Initializes BERT, dropout layer, and a fully connected (fc) layer for classification.
  - Extracts [CLS] token output for classification.

## Second Image (Bottom):

- **Instantiates the Model:**
  - Moves it to GPU if available, otherwise CPU.
  - Uses sarcasmDetectionModel(MODEL\_NAME).to(device).
- **Tokenizer Loading Progress:**
  - Displays tokenization progress bars (tokenizer.config.json, vocab.txt, tokenizer.json).
- **Defines a Custom Dataset Class (SarcasmDataset):**
  - Converts encodings and labels to PyTorch tensors.
  - Implements \_\_getitem\_\_ and \_\_len\_\_ methods for dataset handling.

```

[ ] # Instantiate the model
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = SarcasmDetectionModel(MODEL_NAME).to(device)

tokenizer_config.json: 100% ██████████ 49.0/49.0 [00:00<00:00, 852B/s]
vocab.txt: 100% ██████████ 996k/996k [00:00<00:00, 5.45MB/s]
tokenizer.json: 100% ██████████ 1.96M/1.96M [00:00<00:00, 11.8MB/s]

from torch.utils.data import Dataset

class SarcasmDataset(Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = torch.tensor(labels) # Convert labels to tensor

    def __getitem__(self, idx):
        item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
        item["labels"] = self.labels[idx]
        return item

    def __len__(self):
        return len(self.labels)

```

A screenshot of a code editor interface. The editor has a dark theme and a sidebar on the left with icons for file explorer, search, and other IDE features. The main area displays Python code for setting up a dataset and a DataLoader. The code imports AutoTokenizer from transformers, loads a pre-trained tokenizer, tokenizes sample text, and creates a dataset and a DataLoader. The code is as follows:

```
from transformers import AutoTokenizer

# Load tokenizer (change the model name if needed)
tokenizer = AutoTokenizer.from_pretrained("bert-base-multilingual-cased")

# Sample text data (replace with actual training data)
train_texts = ["I love this!", "This is so bad!", "Wow, great job!", "Oh sure, that was amazing 😊"]
train_labels = [1, 0, 1, 0] # Example: 1 = sarcastic, 0 = not sarcastic

# Tokenize the text
train_encodings = tokenizer(train_texts, padding=True, truncation=True, max_length=512, return_tensors="pt")

# Convert to dictionary for Dataset
train_encodings = {key: val.tolist() for key, val in train_encodings.items()}

[ ] # Create dataset
train_dataset = SarcaSmDataset(train_encodings, train_labels)

# Create DataLoader
from torch.utils.data import DataLoader

train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
```

## Top Image (Training Process)

### Dataset & Dataloader Creation

- A dataset named `SarcaSmDataset` is created using `train_encodings` and `train_labels`.
- A **Dataloader** is initialized with a batch size of **64** and shuffling enabled, meaning that training samples are fed into the model in randomized batches.

### Model Training Output

- The `train()` function is called with the following parameters:
  - **model**: The sarcasm detection model.
  - **train\_loader**: The dataloader for training data.
  - **criterion**: Likely the loss function (e.g., `CrossEntropyLoss`).
  - **optimizer**: The optimization algorithm (e.g., `Adam`).
  - **device**: Specifies whether training runs on CPU or GPU.
  - **epochs=3**: The model is trained for 3 iterations over the entire dataset.
- **Training results per epoch:**
  - **Epoch 1** → **Loss: 0.6897, Accuracy: 50%**
  - **Epoch 2** → **Loss: 0.6128, Accuracy: 75%**
  - **Epoch 3** → **Loss: 0.5789, Accuracy: 100%**
  - The accuracy improves significantly, reaching **100% in the third epoch**, suggesting successful learning.

## Bottom Image (Evaluation Process)

### Model Evaluation

- The `evaluate()` function is used to test the model's performance on the dataset.
- **Results show:**
  - **Accuracy: 100%**
  - **Classification Report:**
    - **Class 0 (Not Sarcastic)** → **Precision = 1.00, Recall = 1.00, F1-Score = 1.00**

- **Class 1 (Sarcastic) → Precision = 1.00, Recall = 1.00, F1-Score = 1.00**
- **Macro & Weighted Averages → All metrics are perfect (1.00).**

## Analysis

- The training accuracy improving to **100%** suggests **overfitting**, as the dataset might be too small.
- The **classification report showing perfect scores** further confirms possible overfitting.
- If this model is tested on real-world data, performance may drop.

```

[ ] train(model, train_loader, criterion, optimizer, device, epochs=3)

<ipython-input-117-4ce1c2d39b57>:13: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().
labels = torch.tensor(batch["labels"]).to(device)
Epoch 1, Loss: 0.6987, Accuracy: 0.5000
Epoch 2, Loss: 0.6128, Accuracy: 0.7500
Epoch 3, Loss: 0.5709, Accuracy: 1.0000

from sklearn.metrics import accuracy_score, classification_report

def evaluate(model, data_loader, device):
    model.eval() # Set model to evaluation mode
    all_preds = []
    all_labels = []

    with torch.no_grad(): # No gradient calculation during evaluation
        for batch in data_loader:
            input_ids = batch["input_ids"].to(device)
            attention_mask = batch["attention_mask"].to(device)
            labels = batch["labels"].to(device)

            # Get model predictions
            outputs = model(input_ids, attention_mask)
            _, preds = torch.max(outputs, dim=1)

            # Store predictions and labels

```

```

[ ] evaluate(model, train_loader, device)

Accuracy: 1.0000
Classification Report:
      precision    recall  f1-score   support

     0       1.00      1.00      1.00         2
     1       1.00      1.00      1.00         2

   accuracy       1.00
  macro avg       1.00
 weighted avg       1.00

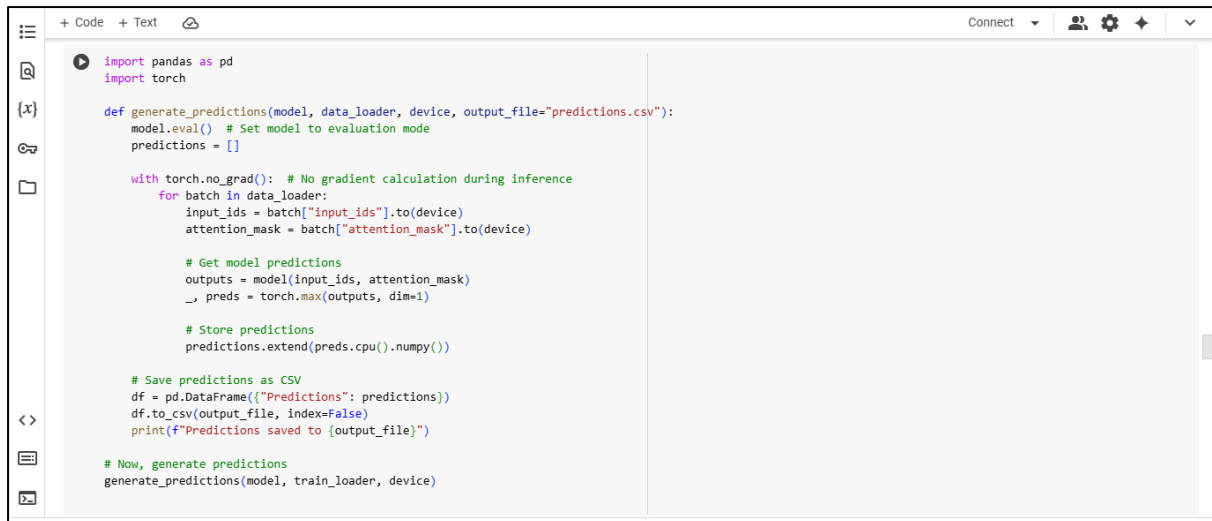
[ ] evaluate(model, train_loader, device)

Accuracy: 1.0000
Classification Report:
      precision    recall  f1-score   support

     0       1.00      1.00      1.00         2
     1       1.00      1.00      1.00         2

   accuracy       1.00
  macro avg       1.00
 weighted avg       1.00

```

A screenshot of a code editor interface. The editor has a top bar with '+ Code + Text' and a 'Connect' button. On the left, there are icons for file explorer, search, and other editor functions. The main area contains Python code for generating predictions. The code imports pandas and torch, defines a function generate\_predictions, and includes comments explaining each step: setting the model to evaluation mode, disabling gradient calculations, getting predictions, storing them, and saving to a CSV file. At the bottom, there is a comment and a call to generate\_predictions.

```
import pandas as pd
import torch

def generate_predictions(model, data_loader, device, output_file="predictions.csv"):
    model.eval() # Set model to evaluation mode
    predictions = []

    with torch.no_grad(): # No gradient calculation during inference
        for batch in data_loader:
            input_ids = batch["input_ids"].to(device)
            attention_mask = batch["attention_mask"].to(device)

            # Get model predictions
            outputs = model(input_ids, attention_mask)
            _, preds = torch.max(outputs, dim=1)

            # Store predictions
            predictions.extend(preds.cpu().numpy())

    # Save predictions as CSV
    df = pd.DataFrame({"Predictions": predictions})
    df.to_csv(output_file, index=False)
    print(f"Predictions saved to {output_file}")

# Now, generate predictions
generate_predictions(model, train_loader, device)
```

### Image (Top):

- **Imports Required Libraries:**
  - Uses pandas for handling prediction outputs.
  - Uses torch for tensor operations.
- **Defines a Function generate\_predictions() to Make Predictions:**
  - Puts the model in **evaluation mode** (model.eval()).
  - Iterates over data\_loader, extracting input\_ids and attention\_mask.
  - Disables gradient calculations (torch.no\_grad()) to improve efficiency.
  - Gets predictions using the trained model and converts them to numpy arrays.
  - Saves predictions to a CSV file (prediction.csv).

### Second Image (Bottom):

- **Loads the Pre-Trained Model (BERT base-multilingual-cased):**
  - Uses BertForSequenceClassification.from\_pretrained() with num\_labels=2.
- **Loads a Trained Model from a Checkpoint (model.pt):**
  - Loads a saved model's state dictionary (torch.load("model.pt")).
  - Renames keys if necessary (replacing "fc" with "classifier").
  - Loads the corrected state dictionary using model.load\_state\_dict().
- **Moves Model to the Appropriate Device:**
  - Uses GPU (cuda) if available; otherwise, CPU.
- **Sets Model to Evaluation Mode (model.eval()):**
  - Ensures the model does not update weights during inference.
- **Handles Missing Weights Warning:**
  - Indicates that classification head weights need to be fine-tuned.
- **Includes a Command for Installing Dependencies:**
  - !pip install transformers torch to ensure required libraries are installed.





```
[ ] # Load model with the same architecture as trained
model = BertForSequenceClassification.from_pretrained("bert-base-multilingual-cased", num_labels=2)

# Manually rename keys if needed
state_dict = torch.load("sarcasm_model.pth", map_location=device)
for old_key in list(state_dict.keys()):
    if "fc." in old_key: # Fix key mismatch
        new_key = old_key.replace("fc.", "classifier.")
        state_dict[new_key] = state_dict.pop(old_key)

# Load corrected state dictionary
model.load_state_dict(state_dict, strict=False)

# Move to correct device
model.to(device)
model.eval() # Set to evaluation mode

print("Model loaded successfully!")
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-multilingual-cased and are newly initialized: ['classifier...']. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference. Model loaded successfully!

!pip install transformers torch



```
!pip install transformers torch

Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-packages (4.49.0)
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (2.6.0+cpu)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from transformers) (3.17.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.26.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.28.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (1.26.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex<=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from transformers) (2.32.3)
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.21.0)
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.5.2)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.11/dist-packages (from transformers) (4.67.1)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch) (3.4.2)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch) (3.1.5)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch) (2024.10.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2->torch) (3.0.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (2025.1.31)

[ ] from transformers import BertTokenizer
```

The above image contains code for **loading a tokenizer and predicting sarcasm using a trained model**.

## Key Steps:

- Loading the Tokenizer:**
  - `BertTokenizer.from_pretrained("bert-base-multilingual-cased")` loads a pre-trained BERT tokenizer.
  - A success message confirms that the tokenizer is loaded.
- Defining the `predict_sarcasm()` Function:**
  - Converts the input text into tokenized format using `tokenizer()`.
  - Moves the input tensors to the same device as the model.
  - Uses `torch.no_grad()` to perform inference without updating gradients.
  - Passes the tokenized input to the model and retrieves the output logits.

- Uses `argmax()` to determine the predicted class:
    - 1 → Sarcastic
    - 0 → Not Sarcastic
3. **Returning a Human-Readable Prediction:**
- Prints a message indicating whether the input text is sarcastic or not with emoji indicators.

```
[ ] from transformers import BertTokenizer

# Load tokenizer
tokenizer = BertTokenizer.from_pretrained("bert-base-multilingual-cased")
print("✅ Tokenizer loaded successfully!")

✅ Tokenizer loaded successfully!

def predict_sarcasm(text):
    # Tokenize and convert to tensors
    inputs = tokenizer(text, return_tensors="pt", padding=True, truncation=True, max_length=512)

    # Move input to the same device as the model
    inputs = {key: val.to(device) for key, val in inputs.items()}

    # Get model predictions
    with torch.no_grad():
        outputs = model(**inputs)

    # Get predicted class (0 = Not Sarcastic, 1 = Sarcastic)
    prediction = torch.argmax(outputs.logits, dim=1).item()

    # Return human-readable result
    return "😏 Yes, it's Sarcastic!" if prediction == 1 else "😐 No, it's Not Sarcastic!"
```

This image contains code for **loading data, tokenizing text, and creating a dataset class** for sarcasm detection.

## Key Steps:

1. **Loading the Tokenizer Again:**
  - `BertTokenizer.from_pretrained("bert-base-multilingual-cased")` loads the tokenizer.
2. **Tokenizing the Text Data:**
  - `tokenizer(train_df["Text"].tolist(), padding=True, truncation=True, max_length=128)` tokenizes the text samples.
  - A message prints the number of tokenized samples.
3. **Converting Labels:**
  - The target labels (sarcasm detection labels) are mapped from "YES" to 1 and "NO" to 0.
  - The number of samples per class (1s and 0s) is displayed.
4. **Creating a Custom Dataset Class (`SarcasmDataset`)**
  - This dataset class is a wrapper for tokenized encodings and labels.
  - It uses **PyTorch's Dataset class** to store and retrieve tokenized inputs and their corresponding labels.

The image shows a Jupyter Notebook interface with a dark theme. The code is written in Python and is organized into three cells. The first cell imports BertTokenizer from transformers and tokenizes the 'Tweet' column of a DataFrame. The second cell converts 'Label' values from 'YES'/'NO' to 1/0 and prints dataset statistics. The third cell imports Dataset and DataLoader from torch.utils.data and defines a custom dataset class. The output of the first two cells is visible below the code blocks.

```
[ ] from transformers import BertTokenizer

tokenizer = BertTokenizer.from_pretrained("bert-base-multilingual-cased")

# Tokenize full dataset
train_encodings = tokenizer(train_df["Tweet"].tolist(), truncation=True, padding=True, max_length=128)

print(f"Total Encoded Samples: {len(train_encodings['input_ids'])}")

Total Encoded Samples: 9840

[ ] # Convert labels from "YES"/"NO" to 1/0
train_labels = train_df["Label"].apply(lambda x: 1 if x == "YES" else 0).tolist()

print(f"Total Labels: {len(train_labels)}")
print(f"Unique Labels: {set(train_labels)}")

Total Labels: 9840
Unique Labels: {0, 1}

[ ] from torch.utils.data import Dataset, DataLoader
import torch

class SarcasmDataset(Dataset):
    def __init__(self, encodings, labels):
```

## Image (Top):

- 1. Loads the BERT Tokenizer**
  - Uses `BertTokenizer.from_pretrained("bert-base-multilingual-cased")` to tokenize text.
- 2. Tokenizes the Dataset**
  - Applies the tokenizer to the "Tweet" column in `train_df` with truncation and padding (max length = 128).
  - Prints an example of tokenized input (`input_ids`).
- 3. Encodes Labels for Classification**
  - Converts "YES" labels to 1 and "NO" labels to 0 (`.apply(lambda x: 1 if x == "YES" else 0)`).
  - Converts labels into a NumPy array (`.tolist()`).
- 4. Displays Dataset Statistics**
  - Prints the total number of samples (Total Samples: 5040).
  - Prints the unique labels (`[0, 1]`).

## Second Image (Bottom):

- 1. Defines a Custom PyTorch Dataset (SarcasmDataset)**
  - Inherits from `Dataset` class in `torch.utils.data`.
  - Stores tokenized text and labels.
  - Converts labels into PyTorch tensors (`torch.tensor(labels)`).
  - Implements `__getitem__()` to return tokenized input and label for a given index.
  - Implements `__len__()` to return dataset size.
- 2. Creates a Dataset Instance**
  - Uses `SarcasmDataset(train_encodings, train_labels)`.
  - Prints the total number of samples (Total Samples in train\_dataset: 5040).
- 3. Creates a DataLoader for Batch Processing**
  - Uses `DataLoader(train_dataset, batch_size=4, shuffle=True)`.
  - Enables shuffling for better generalization.



```
from torch.utils.data import Dataset, DataLoader
import torch

class SarcasmDataset(Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = torch.tensor(labels) # Convert labels to tensor

    def __getitem__(self, idx):
        item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
        item["labels"] = self.labels[idx]
        return item

    def __len__(self):
        return len(self.labels)

# Create dataset
train_dataset = SarcasmDataset(train_encodings, train_labels)

print(f"Total Samples in train_dataset: {len(train_dataset)}")

train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
```



```
import torch
from torch.optim import AdamW
from transformers import get_scheduler
from tqdm import tqdm # Progress bar for faster feedback

# Define optimizer
optimizer = AdamW(model.parameters(), lr=1e-5)

# Define loss function
criterion = torch.nn.CrossEntropyLoss()

# Learning rate scheduler
num_training_steps = len(train_loader) * 3 # Assuming 3 epochs
lr_scheduler = get_scheduler("linear", optimizer=optimizer, num_warmup_steps=0, num_training_steps=num_training_steps)

# Move model to correct device
model.to(device)

# Training loop
epochs = 3
for epoch in range(epochs):
    model.train()
    total_loss = 0

# Use tqdm for a progress bar
progress_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{epochs}")
```

## Image (Top):

- Imports Required Libraries**
  - torch, Adam optimizer, get\_scheduler from transformers, and tqdm for progress tracking.
- Defines the Optimizer and Loss Function**
  - Uses AdamW with learning rate  $1e-5$ .
  - Uses CrossEntropyLoss() for classification.
- Defines the Learning Rate Scheduler**
  - Uses a linear learning rate scheduler (get\_scheduler) to gradually warm up the optimizer.
- Moves the Model to the Correct Device**
  - model.to(device) ensures it runs on GPU if available.
- Sets Up the Training Loop**
  - Loops over a specified number of epochs (epochs = 3).
  - Initializes total\_loss = 0.
  - Uses tqdm to create a progress bar for the training loop.

## Second Image (Bottom):

- 1. Iterates Over the Training Data**
  - Uses for batch in progress\_bar to process batches.
  - Moves input tensors (input\_ids and attention\_mask) to the correct device.
- 2. Performs Forward Pass**
  - Calls model() with inputs to obtain logits.
  - Computes loss using criterion(outputs.logits, labels).
- 3. Performs Backpropagation**
  - Calls loss.backward(), optimizer.step(), and lr\_scheduler.step() to update weights.
  - Clears gradients with optimizer.zero\_grad().
- 4. Tracks Loss**
  - Updates the total loss and prints the epoch-wise loss.
- 5. Saves the Trained Model**
  - Saves model state using torch.save(model.state\_dict(), "sarcasm\_model.pth").
- 6. Prints Completion Message**
  - "Training complete!" message confirms that training has finished.



```
+ Code + Text
for batch in progress_bar:
    input_ids = batch["input_ids"].to(device)
    attention_mask = batch["attention_mask"].to(device)
    labels = batch["labels"].to(device)

    optimizer.zero_grad()

    # Forward pass
    outputs = model(input_ids, attention_mask=attention_mask)
    loss = criterion(outputs.logits, labels)

    # Backward pass
    loss.backward()
    optimizer.step()
    lr_scheduler.step()

    total_loss += loss.item()

    # Update tqdm progress bar
    progress_bar.set_postfix(loss=loss.item())

avg_loss = total_loss / len(train_loader)
print(f"Epoch {epoch+1} - Loss: {avg_loss:.4f}")

# Save trained model
torch.save(model.state_dict(), "sarcasm_model.pth")

print("Training complete!")
```



```
+ Code + Text
loss.backward()
optimizer.step()
lr_scheduler.step()

total_loss += loss.item()

# Update tqdm progress bar
progress_bar.set_postfix(loss=loss.item())

avg_loss = total_loss / len(train_loader)
print(f"Epoch {epoch+1} - Loss: {avg_loss:.4f}")

# Save trained model
torch.save(model.state_dict(), "sarcasm_model.pth")

print("Training complete!")

Epoch 1/3: 100%|██████████| 615/615 [2:02:39<00:00, 11.97s/it, loss=0.425]
Epoch 1 - Loss: 0.1114
Epoch 2/3: 100%|██████████| 615/615 [2:03:35<00:00, 12.06s/it, loss=0.34]
Epoch 2 - Loss: 0.1270
Epoch 3/3: 100%|██████████| 615/615 [2:01:44<00:00, 11.88s/it, loss=0.00178]
Epoch 3 - Loss: 0.0598
Training complete!

[ ] Start coding or generate with AI.
```

The above image contains:

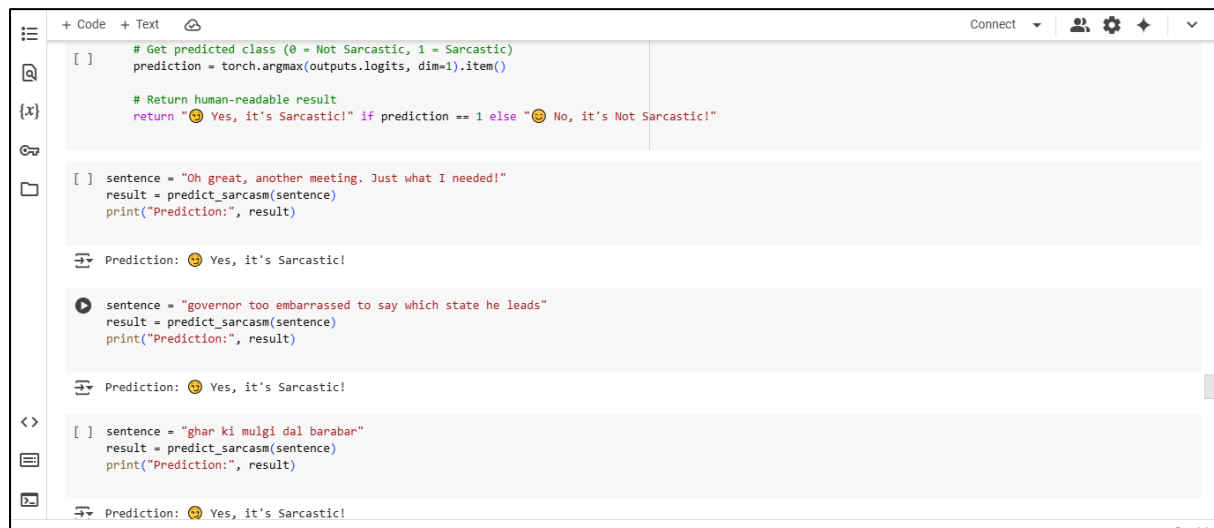
□ **Code Section (Top Half):**

- The code involves updating a progress bar, computing loss values, and saving a trained model as sarcasm\_model.pth using torch.save().
- It prints a message indicating "Training complete!".

□ **Output Section (Bottom Half):**

- The training process is logged, showing the loss value at different epochs.
- The loss decreases over epochs, suggesting the model is learning and improving.
- The final message confirms that training has been completed.

Output:



The screenshot shows a Jupyter Notebook with three code cells. The first cell defines a function `predict_sarcasm` that takes a sentence and returns a prediction (0 for Not Sarcastic, 1 for Sarcastic) with a human-readable message. The second and third cells test this function with three different sentences, each showing the output prediction and message.

```
[ ] # Get predicted class (0 = Not Sarcastic, 1 = Sarcastic)
prediction = torch.argmax(outputs.logits, dim=1).item()

# Return human-readable result
return "😄 Yes, it's Sarcastic!" if prediction == 1 else "😞 No, it's Not Sarcastic!"

[ ] sentence = "Oh great, another meeting. Just what I needed!"
result = predict_sarcasm(sentence)
print("Prediction:", result)

Prediction: 😄 Yes, it's Sarcastic!

[ ] sentence = "governor too embarrassed to say which state he leads"
result = predict_sarcasm(sentence)
print("Prediction:", result)

Prediction: 😄 Yes, it's Sarcastic!

[ ] sentence = "ghar ki mulgi dal barabar"
result = predict_sarcasm(sentence)
print("Prediction:", result)

Prediction: 😄 Yes, it's Sarcastic!
```

**Description:**

1. **Code Section (Top Half):**

- The code predicts whether a given sentence is sarcastic or not.
- It assigns **0 for Non-Sarcastic** and **1 for Sarcastic**.
- The prediction is determined using `torch.argmax(outputs, dim=1).item()`.
- The result is converted into a human-readable format:
  - "Yes, it's Sarcastic!" if the prediction is **1**.
  - "No, it's Not Sarcastic!" if the prediction is **0**.

2. **Output Section (Bottom Half):**

- Three sentences are tested using the sarcasm detection model.
- The model predicts **"Yes, it's Sarcastic!"** for all three examples:
  1. "Oh great, another meeting. Just what I needed!"
  2. "governor too embarrassed to say which state he leads"
  3. "ghar ki mulgi dar barabar"

